

Computer Vision Libraries

Computer Vision Libraries

- OpenCV Installation

- Introduction

- Basic Image Container

- Image Storage


- Space Domain Techniques

- Point Processing Techniques

- Kernels

- Edges

- Video



OpenCV Installation

1 .- Install the following required packages

```
$ sudo apt-get install build-essential
$ sudo apt-get install cmake git libgtk2.0-dev pkg-config libavcodec-dev libavformat-dev libswscale-dev
$ sudo apt-get install python-dev python-numpy libtbb2 libtbb-dev libjpeg-dev libpng-dev libtiff-dev
libjasper-dev libdc1394-22-dev
```

2 .- Getting OpenCV Source Code

```
$ mkdir opencv
$ cd opencv
$ git clone https://github.com/opencv/opencv.git
```

3 .- Building OpenCV from Source Using CMake

```
$ cd opencv
$ mkdir build
$ cd build
```

4 .- Configuring, Building and Install

```
$ cmake -D CMAKE_BUILD_TYPE=Release
        -D CMAKE_INSTALL_PREFIX=/usr/local ..
$ make
$ sudo make install
```

NOTE: You may need to copy by hand the commands (ctrl+c may not work)

Introduction

Load and Display an Image using OpenCV

```
#include <opencv2/core.hpp>
#include <opencv2/imgcodecs.hpp>
#include <opencv2/highgui.hpp>
#include <stdio.h>

using namespace cv;

int main(int argc, char** argv )
{
    if ( argc != 2 )
    {
        printf("usage: DisplayImage <Image_Path>\n");
        return -1;
    }
    Mat image;
    image = imread( argv[1], IMREAD_COLOR );
    if ( !image.data )
    {
        printf("No image data \n");
        return -1;
    }
    namedWindow("Display Image", WINDOW_AUTOSIZE );
    imshow("Display Image", image);
    waitKey(0);
    return 0;
}
```

Create a CMake file

```
cmake_minimum_required(VERSION 2.8)
project( DisplayImage )
find_package( OpenCV REQUIRED )
include_directories( ${OpenCV_INCLUDE_DIRS} )
add_executable( DisplayImage DisplayImage.cpp )
target_link_libraries( DisplayImage ${OpenCV_LIBS} )
```

Generate the executable

```
$ cd <DisplayImage_directory>
$ cmake .
$ make
```

Run it giving an image location as an argument

```
$ ./DisplayImage lena.jpg
```

Introduction

- ▶ **core** section, as here are defined the basic building blocks of the library
- ▶ **highgui** module, as this contains the functions for input and output operations

```
#include <opencv2/core.hpp>
#include <opencv2/imgcodecs.hpp>
#include <opencv2/highgui.hpp>
```

OpenCV has its own namespace: `cv`

```
using namespace cv;
```

We start up assuring that we acquire a valid image name argument from the command line:

```
if ( argc != 2 )
{
    printf("usage: DisplayImage <Image_Path>\n");
    return -1;
}
```

Then create a `Mat` object that will store the data of the loaded image.

```
Mat image;
```

Now we call the **`cv::imread`** function which loads the image name specified by the first argument (`argv[1]`). The second argument specifies the format in what we want the image. This may be:

```
image = imread( argv[1], IMREAD_COLOR );
```

After checking that the image data was loaded correctly, we want to display our image, so we create an OpenCV window using the **`cv::namedWindow`** function and finally, to update the content of the OpenCV window with a new image use the **`cv::imshow`** function

```
namedWindow("Display Image", WINDOW_AUTOSIZE );
imshow("Display Image", image);
```

Introduction

Modify and Save an Image using OpenCV

```
#include <opencv2/opencv.hpp>

using namespace cv;

int main( int argc, char** argv )
{
    char* imageName = argv[1];
    Mat image;
    image = imread( imageName, IMREAD_COLOR );
    if( argc != 2 || !image.data )
    {
        printf( " No image data \n " );
        return -1;
    }
    Mat gray_image;
    cvtColor( image, gray_image, COLOR_BGR2GRAY );
    imwrite( "Gray_Image.jpg", gray_image );
    namedWindow( imageName, WINDOW_AUTOSIZE );
    namedWindow( "Gray image", WINDOW_AUTOSIZE );
    imshow( imageName, image );
    imshow( "Gray image", gray_image );
    waitKey(0);
    return 0;
}
```

Now we are going to convert our image from BGR to Grayscale format. OpenCV has a really nice function to do this kind of transformations:

```
cvtColor( image, gray_image, COLOR_BGR2GRAY );
```


To save it, we will use a function analogous to cv::imread : **cv::imwrite**

```
imwrite( "../Gray_Image.jpg", gray_image );
```

Add the **waitKey(0)** function call for the program to wait forever for an user key press.

```
waitKey(0);
```

Basic Image Container



The mirror of the car is nothing more than a matrix containing all the intensity values of the pixel points.

Mat

Mat is a class with two data parts:

- ▶ **The matrix header** containing information such as the size of the matrix, the method used for storing, at which address is the matrix stored, and so on
- ▶ **Pointer to the matrix** containing the pixel values

Mat has the following properties:

- ▶ Output image allocation for OpenCV functions is automatic (unless specified otherwise).
- ▶ You do not need to think about memory management with OpenCV's C++ interface.

- ▶ The assignment operator and the copy constructor only copies the header.
- ▶ The underlying matrix of an image may be copied using the **cv::Mat::clone()** and **cv::Mat::copyTo()** functions.

```
Mat matA, matC;                                // creates just the header parts
matA = imread(argv[1], IMREAD_COLOR); // allocate matrix
Mat matB(matA);                                // Use the copy constructor
matC = matA;                                   // Assignment operator
Mat matD(matA, Rect(10, 10, 100, 100) ); // using a rectangle
Mat matE = matA(Range::all(), Range(1,3)); // using row and column
Mat matF = matA.clone();
Mat matG;
matA.copyTo(matG);
```

Mat

Create a Mat object in multiple ways

- ▶ **cv::Mat::Mat** Constructor

```
Mat matM(2,2, CV_8UC3, Scalar(0,0,255));
```

- ▶ Use C/C++ arrays and initialize via constructor

```
int varsz[3] = {2,2,2};  
Mat test(3,varsz, CV_8UC(1), Scalar::all(0));
```

- ▶ **cv::Mat::create** function

```
Mat.create(4,4, CV_8UC(2));
```

- ▶ For small matrices you may use comma separated

```
Mat matC = (Mat_<double>(3,3) << 0, -1, 0, -1, 5, -1, 0, -1, 0);
```

Output formatting

- ▶ Default

```
cout << "matR (default) = " << endl << matR;
```

- ▶ Comma separated values (CSV)

```
cout << "matR (csv)= " << endl << format(matR, Formatter::FMT_CSV );
```

- ▶ 2D Point

```
Point2f poP(5, 1);  
cout << "Point (2D) = " << poP;
```

- ▶ 3D Point

```
Point3f P3f(2, 6, 7);  
cout << "Point (3D) = " << P3f;
```


Image Storage

The size of the matrix depends on the color system used. More accurately, it depends from the number of channels used. In case of a gray scale image we have something like:

	Column 0	Column 1	Column ...	Column m
Row 0	0,0	0,1	...	0, m
Row 1	1,0	1,1	...	1, m
Row,0	...,1, m
Row n	n,0	n,1	n,...	n, m

For multichannel images the columns contain as many sub columns as the number of channels. For example in case of an BGR color system:

	Column 0			Column 1			Column ...			Column m		
Row 0	0,0	0,0	0,0	0,1	0,1	0,1	0, m	0, m	0, m
Row 1	1,0	1,0	1,0	1,1	1,1	1,1	1, m	1, m	1, m
Row,0	...,0	...,0	...,1	...,1	...,1, m	..., m	..., m
Row n	n,0	n,0	n,0	n,1	n,1	n,1	n,...	n,...	n,...	n, m	n, m	n, m

The efficient method

The classic C style operator[] (pointer) access:

```
Mat& ScanImageAndReduceC(Mat& matI, const uchar* const table)
{
    CV_Assert(matI.depth() == CV_8U); //only char
    int channels = matI.channels();
    int nRows = matI.rows;
    int nCols = matI.cols * channels;
    if (matI.isContinuous())
    {
        nCols /= nRows;
        nRows = 1;
    }
    int idx,jdx;
    uchar* ptr;
    for( idx = 0; idx < nRows; ++idx)
    {
        ptr = matI.ptr<uchar>(idx);
        for ( jdx = 0; jdx < nCols; ++jdx)
        {
            ptr[jdx] = table[ptr[jdx]];
        }
    }
    return matI;
}
```

- ▶ Basically just acquire a pointer to the start of each row and go through it until it ends
- ▶ In the special case that the matrix is stored in a continuous manner we only need to request the pointer a single time and go all the way to the end
- ▶ to look out for color images: we have three channels so we need to pass through three times more items in each row

The iterator (safe) method

```
Mat& ScanImageAndReduceIterator(Mat& matI, const uchar* const table)
{
    CV_Assert(matI.depth() == CV_8U);    //only char
    const int channels = matI.channels();
    switch(channels)
    {
        case 1:
        {
            MatIterator_<uchar> it, end;
            for( it = matI.begin<uchar>(), end = matI.end<uchar>(); it != end; ++it)
                *it = table[*it];
            break;
        }
        case 3:
        {
            MatIterator_<Vec3b> it, end;
            for( it = matI.begin<Vec3b>(), end = matI.end<Vec3b>(); it != end; ++it)
            {
                (*it)[0] = table[(int)(*it)[0]];
                (*it)[1] = table[(int)(*it)[1]];
                (*it)[2] = table[(int)(*it)[2]];
            }
        }
    }
    return matI;
}
```

- ▶ The iterator method is considered a safer way as it takes over these tasks from the user
- ▶ Ask the begin and the end of the image matrix and then just increase the begin iterator until you reach the end
- ▶ In case of color images we have three uchar items per column (Vec3b)

```
Vec3b intensity = img.at<Vec3b>(y, x);
uchar blue = intensity.val[0];
uchar green = intensity.val[1];
uchar red = intensity.val[2];
```

Performance Difference


To make the differences more clear, a quite large (2560 X 1600) image is used. The performance presented here are for color images. For a more accurate value I've averaged the value I got from the call of the function for hundred times.

Method	Time
Efficient	79.4717 milliseconds
Iterator	83.7201 milliseconds

Space Domain Techniques - Histogram

- ▶ Histograms are collected counts of data organized into a set of predefined bins
- ▶ The data collected can be whatever feature you find useful to describe your image
- ▶ Since we know that the range of information value, we can segment our range in subparts (**bins**)

Imagine that a Matrix contains information of an image (i.e. intensity in the range 0-255):



OpenCV implements the function **cv::calcHist**, which calculates the histogram of a set of arrays (usually images or image planes). It can operate with up to 32 dimensions.

Space Domain Techniques - Histogram

```
#include "opencv2/highgui.hpp"
#include "opencv2/imgcodecs.hpp"
#include "opencv2/imgproc.hpp"
#include <iostream>

int main(int argc, char** argv)
{
    Mat src; //... Image readed
    vector<Mat> bgr_planes;
    split( src, bgr_planes );
    int histSize = 256;
    float range[] = { 0, 256 }; //the upper boundary is exclusive
    const float* histRange = { range };
    bool uniform = true, accumulate = false;
    Mat b_hist, g_hist, r_hist;
    calcHist( &bgr_planes[0], 1, 0, Mat(), b_hist, 1, &histSize,
             &histRange, uniform, accumulate );
    calcHist( &bgr_planes[1], 1, 0, Mat(), g_hist, 1, &histSize,
             &histRange, uniform, accumulate );
    calcHist( &bgr_planes[2], 1, 0, Mat(), r_hist, 1, &histSize,
             &histRange, uniform, accumulate );
    int hist_w = 512, hist_h = 400;
    int bin_w = cvRound( (double) hist_w/histSize );
    Mat histImage( hist_h, hist_w, CV_8UC3, Scalar( 0,0,0) );
    normalize(b_hist, b_hist, 0, histImage.rows, NORM_MINMAX, -1, Mat() );
    normalize(g_hist, g_hist, 0, histImage.rows, NORM_MINMAX, -1, Mat() );
    normalize(r_hist, r_hist, 0, histImage.rows, NORM_MINMAX, -1, Mat() );
```

```
for( int i = 1; i < histSize; i++ )
{
    line( histImage,
          Point( bin_w*(i-1), hist_h - cvRound(b_hist.at<float>(i-1)) ),
          Point( bin_w*(i), hist_h - cvRound(b_hist.at<float>(i)) ),
          Scalar( 255, 0, 0 ), 2, 8, 0 );
    line( histImage,
          Point( bin_w*(i-1), hist_h - cvRound(g_hist.at<float>(i-1)) ),
          Point( bin_w*(i), hist_h - cvRound(g_hist.at<float>(i)) ),
          Scalar( 0, 255, 0 ), 2, 8, 0 );
    line( histImage,
          Point( bin_w*(i-1), hist_h - cvRound(r_hist.at<float>(i-1)) ),
          Point( bin_w*(i), hist_h - cvRound(r_hist.at<float>(i)) ),
          Scalar( 0, 0, 255 ), 2, 8, 0 );
}
imshow("Source image", src );
imshow("calcHist Demo", histImage );
waitKey();
return 0;
}
```

Space Domain Techniques - Histogram

Calculate the histograms by using the OpenCV function **cv::calcHist**

```
Mat b_hist, g_hist, r_hist;  
calcHist( &bgr_planes[0], 1, 0, Mat(), b_hist, 1, &histSize, &histRange, uniform, accumulate );  
calcHist( &bgr_planes[1], 1, 0, Mat(), g_hist, 1, &histSize, &histRange, uniform, accumulate );  
calcHist( &bgr_planes[2], 1, 0, Mat(), r_hist, 1, &histSize, &histRange, uniform, accumulate );
```

- ▶ **&bgr_planes[0]** The source array(s)
- ▶ **1** The number of source arrays (in this case we are using 1. We can enter here also a list of arrays)
- ▶ **0** The channel (dim) to be measured. In this case it is just the intensity (each array is single-channel) so we just write 0.
- ▶ **Mat()** A mask to be used on the source array (zeros indicating pixels to be ignored). If not defined it is not used
- ▶ **b_hist** The Mat object where the histogram will be stored
- ▶ **1** The histogram dimensionality.
- ▶ **histSize** The number of bins per each used dimension
- ▶ **histRange** The range of values to be measured per each dimension
- ▶ **uniform and accumulate** The bin sizes are the same and the histogram is cleared at the beginning.

Space Domain Techniques - Histogram

cv::normalize the histogram so its values fall in the range indicated by the parameters entered

```
normalize(b_hist, b_hist, 0, histImage.rows, NORM_MINMAX, -1, Mat() );  
normalize(g_hist, g_hist, 0, histImage.rows, NORM_MINMAX, -1, Mat() );  
normalize(r_hist, r_hist, 0, histImage.rows, NORM_MINMAX, -1, Mat() );
```


- ▶ **b_hist** Input array
- ▶ **b_hist** Output normalized array (can be the same)
- ▶ **0** and **histImage.rows** For this example, they are the lower and upper limits to normalize the values of **r_hist**
- ▶ **NORM_MINMAX** Argument that indicates the type of normalization (it adjusts the values between the two limits set before)
- ▶ **b_hist** The Mat object where the histogram will be stored
- ▶ **-1** Implies that the output normalized array will be the same type as the input
- ▶ **Mat()** Optional mask

Space Domain Techniques - Histogram

Using as input argument an image like this one




Produces the following histogram



Space Domain Techniques - Histogram Equalization

- ▶ It is a method that improves the contrast in an image, in order to stretch out the intensity range
- ▶ You can see that the pixels seem clustered around the middle of the available range of intensities. What **Histogram Equalization** does is to stretch out this range. The green circles indicate the underpopulated intensities. After applying the equalization, we get an histogram like the figure in the center. The resulting image is shown in the picture at right.



Space Domain Techniques - Histogram Equalization

```
#include "opencv2/imgcodecs.hpp"
#include "opencv2/highgui.hpp"
#include "opencv2/imgproc.hpp"
#include <iostream>
using namespace cv;
using namespace std;
int main( int argc, char** argv )
{
    Mat src; //... Image readed
    cvtColor( src, src, COLOR_BGR2GRAY );
    Mat dst;
    equalizeHist( src, dst );
    imshow( "Source image", src );
    imshow( "Equalized Image", dst );
    waitKey();
    return 0;
}
```

► Convert it to grayscale

```
cvtColor( src, src, COLOR_BGR2GRAY );
```

► Apply histogram equalization with the function **cv::equalizeHist**



```
Mat dst;
equalizeHist( src, dst );
```

Space Domain Techniques - Histogram Equalization

Source Image




Equalized Image



Point Processing Techniques - Threshold

- ▶ The simplest segmentation method
- ▶ Separate out regions of an image corresponding to objects which we want to analyze
- ▶ To differentiate the pixels we are interested in from the rest
- ▶ Once we have separated properly the important pixels, we can set them with a determined value to identify them



Point Processing Techniques - Threshold

```
#include "opencv2/highgui.hpp"
#include "opencv2/imgproc.hpp"
#include <stdlib.h>
#include <stdio.h>

// Global variables

int threshold_value = 0;
int threshold_type = 3;;
int const max_value = 255;
int const max_type = 4;
int const max_BINARY_value = 255;

Mat src, src_gray, dst;
char* window_name = "Threshold Demo";

char* trackbar_type = "Type: \n 0: Binary \n 1: Binary Inverted \n"
    "2: Truncate \n 3: To Zero \n 4: To Zero Inverted";
char* trackbar_value = "Value";

/// Function headers
void Threshold_Demo( int, void* );

int main( int argc, char** argv )
{
    /// Load an image
    src = imread( argv[1], 1 );
```

```
cvtColor( src, src_gray, CV_BGR2GRAY );
namedWindow( window_name, CV_WINDOW_AUTOSIZE );
createTrackbar( trackbar_type,
                window_name, &threshold_type,
                max_type, Threshold_Demo );
createTrackbar( trackbar_value,
                window_name, &threshold_value,
                max_value, Threshold_Demo );
Threshold_Demo( 0, 0 );
while(true)
{
    int c;
    c = waitKey( 20 );
    if( (char)c == 27 )
    { break; }
}
}

void Threshold_Demo( int, void* )
{
    /* 0: Binary
       1: Binary Inverted
       2: Threshold Truncated
       3: Threshold to Zero
       4: Threshold to Zero Inverted
    */
    threshold( src_gray, dst, threshold_value, max_BINARY_value,
               threshold_type );
    imshow( window_name, dst );
}
```

Point Processing Techniques - Threshold

- ▶ Create 2 trackbars for the user to enter user input

```
createTrackbar( trackbar_type, window_name, &threshold_type,  
               max_type, Threshold_Demo );  
createTrackbar( trackbar_value, window_name, &threshold_value,  
               max_value, Threshold_Demo );
```

- ▶ Wait until the user enters the threshold value, the type of thresholding (or until the program exits)
- ▶ Whenever the user changes the value of any of the Trackbars, the function Threshold_Demo is called:

```
void Threshold_Demo( int, void* )
```

The function **threshold**:

```
threshold( src_gray, dst, threshold_value, max_BINARY_value,  
           threshold_type );
```

- ▶ **src_gray** Our input image
- ▶ **dst** Destination (output) image
- ▶ **threshold_value** The thresh value with respect to which the thresholding operation is made
- ▶ **max_BINARY_value** The value used with the Binary thresholding operations (to set the chosen pixels)
- ▶ **threshold_type** One of the 5 thresholding operations. They are listed in the comment section of the function above

Point Processing Techniques - Threshold



512 x 512 pixels 549.7 KB 100%

5 / 8



Point Processing Techniques - Log

```
#include "opencv2/highgui.hpp"
#include "opencv2/imgproc.hpp"

int main ()
{
    cv::Mat binary = cv::imread("path",0);

    cv::Mat fg;
    binary.convertTo(fg,CV_32F);
    fg = fg + 1;
    cv::log(fg,fg);
    cv::convertScaleAbs(fg,fg);
    cv::normalize(fg,fg,0,255,cv::NORM_MINMAX);
    cv::imshow("a",fg);
    cv::waitKey(0);
}
```



Kernels

The idea is that we recalculate each pixels value in an image according to a mask matrix (also known as kernel). This mask holds values that will adjust how much influence neighboring pixels (and the current pixel) have on the new pixel value. From a mathematical point of view we make a weighted average, with our specified values.

Consider, for example, the issue of an image contrast enhancement method. Basically we want to apply for every pixel of the image the following formula:

$$I(i,j) = 5 * I(i,j) - [I(i-1,j) + I(i+1,j) + I(i,j-1) + I(i,j+1)]$$
$$\iff I(i,j) * M, \text{ where } M = \begin{matrix} & \begin{matrix} i \backslash j & -1 & 0 & +1 \end{matrix} \\ \begin{matrix} -1 \\ 0 \\ +1 \end{matrix} & \begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix} \end{matrix}$$

Smoothing Images

Smoothing, also called blurring, is a simple and frequently used image processing operation.

There are many reasons for smoothing, for example, to reduce noise.

To perform a smoothing operation we will apply a filter to our image. The most common type of filters are linear, in which an output pixel's value (i.e. $g(i, j)$) is determined as a weighted sum of input pixel values (i.e. $f(i + k, j + l)$) :

$$g(i, j) = \sum_{k, l} f(i + k, j + l)h(k, l)$$

$h(k, l)$ is called the kernel, which is nothing more than the coefficients of the filter.

- ▶ Normalized Box Filter
- ▶ Gaussian Filter
- ▶ Median Filter
- ▶ Bilateral Filter

Borders

One problem that naturally arises is how to handle the boundaries. How can we convolve them if the evaluated points are at the edge of the image?

What most of OpenCV functions do is to copy a given image onto another slightly larger image and then automatically pads the boundary (by any of the methods explained in the sample code just below). This way, the convolution can be performed over the needed pixels without problems (the extra padding is cut after the operation is done).



Laplacian

The Laplacian operator is defined by:

$$\mathcal{L}aplace(f) = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

The Laplacian operator is implemented in OpenCV by the function `Laplacian`.



Video



Open a camera with

```
cv::VideoCapture input_video;  
input_video.open(camera_id);
```

Check if it is open

```
if(!input_video.isOpened())
```

Read a frame with

```
input_video.read(frame_input);
```

At the end, release the camera

```
input_video.release();
```